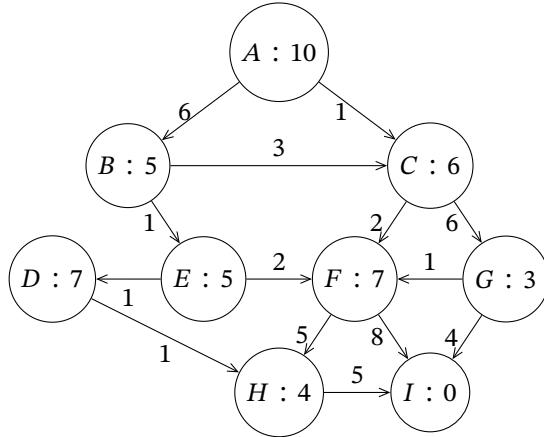


## Assignment3 – Search

### Problem 3.1 (Search Algorithms)

Consider the following *directed graph*:



Every node is *labeled* with  $n : h(n)$  where  $n$  is the identifier of the node and  $h(n)$  is the *heuristic* for estimating the *cost* from  $n$  to a *goal node*. Every edge is *labeled* with its actual *cost*.

1. Assume that  $I$  is the *goal node*. Argue whether or not the *heuristic* is *admissible*.

Now assume you have already *expanded* the node  $A$ . List the **next 4 nodes (i.e., excluding A)** that will be *expanded* using the respective *algorithm*. If there is a tie, break it using *alphabetical order*.

2. *depth-first search*
3. *breadth-first search*
4. *uniform-cost search*
5. *greedy-search*
6.  $A^*$ -*search*

### Problem 3.2 (Formally Modeling a Search Problem)

Consider the Towers of Hanoi for 7 disks initially stacked on peg A.

Is this problem *deterministic*? Is it *fully observable*?

*Formally model* it as a *search problem* in the sense of the *mathematical definition* from the slides. *Explain* how your *mathematical definition* models the problem.

Note that the *formal model* only defines the problem — we are not looking for solutions here.

Note that modeling the problem corresponds to defining it in a *programming language*, except that we use *mathematics* instead of a *programming language*. Then *explaining* the model corresponds to documenting your *implementation*.

### Problem 3.3 (Heuristic Searches)

Consider the graph of Romanian cities with edges labeled with costs  $c(m, n)$  of going from  $m$  to  $n$ .  $c(m, n)$  is always bigger than the *straight-line distance* from  $m$  to  $n$ .  $c(m, n)$  is *infinite* if there is no edge.

Our search algorithm keeps:

- a list  $E$  of expanded nodes  $n$  together with the cost  $g(n)$  of the cheapest path to  $n$  found so far,
- a fringe  $F$  containing the unexpanded neighbors of expanded nodes.

We want to find a cheap path from Lugoj to Bucharest. Initially,  $E$  is *empty*, and  $F$  contains only Lugoj. We *terminate* if  $E$  contains Bucharest.

*Expansion* of a node  $n$  in  $F$  moves it from  $F$  to  $E$  and adds to  $F$  every neighbor of  $n$  that is not already in  $E$  or  $F$ . We obtain  $g(n)$  by minimizing  $g(e) + c(e, n)$  over expanded nodes  $e$ .

As a *heuristic*  $h(n)$ , we use the *straight-line distance* from  $n$  to Bucharest as given by the table in the *lectures*.

Explain how the following algorithms choose which node to expand next:

1. greedy search with heuristic  $h$
2. A\* search with path cost  $g$  and heuristic  $h$
3. Explain what  $h^*$  is here and why  $h$  is *admissible*.
4. For each search, give the order in which nodes are expanded.

(You only have to give the nodes to get the full score. But to get partial credit in case you're wrong, you may want to include for each step all nodes in the fringe and their cost.)

### Problem 3.4 (Heuristics)

Consider heuristic search with heuristic  $h$ .

1. Briefly explain what is the same and what is different between A\* search and greedy search regarding the decision which node to expand next.
2. Is the constant function  $h(n) = 0$  an admissible heuristic for A\* search?

### Problem 3.5 (Tree Search in ProLog)

Implement the following tree search algorithms in Prolog:

1. BFS
2. DFS
3. Iterative Deepening (with step size 1)

Remarks:

- In the *lectures*, we talked about expanding *nodes*. That is relevant in many AI *applications* where the *tree* is not built yet (and maybe even too big to hold in *memory*), such as *game trees* in *move-based games* or *decision trees* of *agents interacting* with an *environment*. In those cases, when visiting a *node*, we have to *expand* it, i.e., *compute* what its *children* are.

In this problem, we work with smaller *trees* where the *search algorithm* receives the fully *expanded tree* as *input*. The *algorithm* must still visit every *node* and perform some operation on it — the *search algorithm* determines in which *order* the *nodes* are visited.

In our case, the operation will be to *write out* the label of the *node*.

- In the *lectures*, we worked with *goal nodes*, where the *search* stops when a *goal node* is found. Here we do something simpler: we *visit all the nodes and operate on each one* without using a *goal state*. (Having a *goal state* is then just the special case where the operation is to *test* the *node* and possibly stop.)

*Concretely*, your submission **must** be a single *Prolog file* that extends the following *implementation*:

```
% tree(V,TS) represents a tree.
% V must be a string - the label/value/data V of the root node
% TS must be a list of trees - the children/subtrees of the root node
% In particular, a leaf node is a tree with the empty list of children
istree(tree(V,TS)) :- string(V), istreelist(TS).

% istreelist(TS) holds if TS is a list of trees.
% Note that the children are a list not a set, i.e., they are ordered.
istreelist([]).
istreelist([T|TS]) :- istree(T), istreelist(TS).

% The following predicates define search algorithms that take a tree T
% and visit every node each time writing out the line D:L where
% * D is the depth (starting at 0 for the root)
% * L is the label

% dfs(T) visits every node in depth-first order
dfs(T) :- ????
% bfs(T) visits every node in breadth-first order
bfs(T) :- ????
% itd(T) :- visits every node in iterative deepening order
itd(T) :- ????
```

Here “must” means you can define any *number* of additional *predicates*. But the *predicates* specified above must exist and must have that *arity* and must work correctly on any *input* T that satisfies *istree*(T). “Working correctly” means the *predicates* must write out exactly what is specified, e.g.,

```
0:A
1:B
```

for the *depth-first search* of the *tree* *tree*("A", [*tree*("B", [])]).