


Assignment3 – Search

Problem 3.1 (Search Algorithms)

Consider the following directed graph:



Every **node** is **labeled** with $n : h(n)$ where n is the identifier of the **node** and $h(n)$ is the **heuristic** for estimating the **cost** from n to a **goal node**. Every **edge** is **labeled** with its actual **cost**.

1. Assume that I is the **goal node**. Argue whether or not the **heuristic** is **admissible**.

Solution: It is not **admissible**: The **cost** from D to the **goal** is $1 + 5 = 6 < 7 = h(D)$, and a **heuristic** must not overestimate that **cost**.

Now assume you have already **expanded** the **node** A . List the **next 4 nodes (i.e., excluding A)** that will be **expanded** using the respective **algorithm**. If there is a tie, break it using **alphabetical order**.

As presented in the notes, the same state may be inserted into the fringe multiple times if it is encountered along multiple paths. The second insertion may happen *before* or *after* the first one has been expanded.

- before: Here both nodes are in the fringe at the same time, and a practical implementation can easily remove duplicates from the fringe as follows:
 - DFS: keep the later node
 - BFS: keep the earlier node
 - UCS, A^* : keep the node with the lower cost
 - greedy: keep either node (they're identical)

This can speed up search and possibly avoid some cycles.

- after: Here the fringe does not contain the same state twice, and an implementation would have to perform a more sophisticated analysis to check if a state has been expanded previously. This would avoid all redundant duplicate expansion and cycles. However, in the case of UCS and A^* , if the second occurrence of the state has a lower cost, the state would still have to be expanded a second time.

In the sequel, we assume that the before-check happens but not the after-check.

2. depth-first search

Solution: B, C, F, H

3. breadth-first search

Solution: B, C, E, F (Without the before-check, we would get B, C, C, E .)

4. uniform-cost search

Solution: C, F, B, E

5. greedy-search

Solution: B, E, C, G

6. A^* -search

Solution: C, F, G, B

Problem 3.2 (Formally Modeling a Search Problem)

Consider the Towers of Hanoi for 7 disks initially stacked on peg A.

Is this problem **deterministic**? Is it **fully observable**?

Formally model it as a **search problem** in the sense of the **mathematical definition** from the slides. Explain how your **mathematical definition** models the problem.

Note that the **formal model** only defines the problem — we are not looking for solutions here.

Note that modeling the problem corresponds to defining it in a [programming language](#), except that we use [mathematics](#) instead of a [programming language](#). Then [explaining](#) the model corresponds to documenting your [implementation](#).

Solution: We need to give $(S, \mathcal{A}, T, I, G)$.

Because the problem is [deterministic](#), we know $\#(T(a, s)) \leq 1$. Because the problem is [fully observable](#), we know $\#(I) = 1$.

Let $D = \{1, 2, 3, 4, 5, 6, 7\}$ (the [set](#) of disks) and $P = \{A, B, C\}$ (the [set](#) of pegs). We put:

- $S = D \rightarrow P$, i.e., a [state](#) s is a [function](#) from disks to pegs.

Explanation: In [state](#) s , the [value](#) $s(d)$ is the peg that disk d is on. Because disks must always be [ordered](#) by size, we do not have to explicitly store the [order](#) in which the disks sit on the pegs.

- $\mathcal{A} = \{(A, B), (B, A), (A, C), (C, A), (B, C), (C, B)\}$, i.e., an [action](#) a is a [pair](#) of different pegs.

Explanation: (p, q) represents the [action](#) of moving the top disk of peg p to peg q .

- For $s \in S$ and $p \in P$, we abbreviate as $\text{top}(s, p)$ the smallest $d \in D$ such that $s(d) = p$.

Explanation: $\text{top}(s, p)$ is the top (smallest) disk on peg p in [state](#) s .

Then $T : \mathcal{A} \times S \rightarrow \mathcal{P}(S)$ is defined as follows:

- If $\text{top}(s, q) > \text{top}(s, p)$, we put $T((p, q), s) = \{s'\}$ where $s' : D \rightarrow P$ is given by
 - * $s'(d) = q$ if $d = \text{top}(s, p)$
 - * $s'(d) = s(d)$ for all other [values](#) of d
- otherwise, $T((p, q), s) = \emptyset$

Explanation: In [state](#) s , if the top disk of q is bigger than the top disk of p , the [action](#) (p, q) is [applicable](#), and the [successor states](#) s' of s after applying (p, q) is the same as s except that the top (smallest) disk on peg p is now on peg q .

- $I = \{i\}$ where the [state](#) i is given by $i(d) = A$ for all $d \in D$.

Explanation: Initially, all disks are on peg A .

- $G = \{g\}$ where the [state](#) g is given by $g(d) = B$ for all $d \in D$.

Explanation: There is only one [goal state](#), described by all disks being on peg B .

Note that there are many different correct solutions to this problem. In particular, you can use different definitions for S (i.e., model the [state space](#) differently), in which case everything else in the model will be different, too. Often a good model for the [state space](#) can be recognized by how straightforward it is to define the rest of the model formal.

Even if you used a different **state space**, a good self-study exercise is to check that the above (a) is indeed a **search problem** and (b) correctly models the Towers of Hanoi. Continuing the above analogy to **programming languages**, (a) corresponds to **compiling/type-checking** your **implementation** and (b) to checking that your **implementation** is correct.

Problem 3.3 (Heuristic Searches)

Consider the **graph** of Romanian cities with **edges labeled** with costs $c(m, n)$ of going from m to n . $c(m, n)$ is always bigger than the **straight-line distance** from m to n . $c(m, n)$ is **infinite** if there is no **edge**.

Our **search algorithm** keeps:

- a **list** E of **expanded nodes** n together with the **cost** $g(n)$ of the cheapest **path** to n found so far,
- a **fringe** F containing the **unexpanded neighbors** of **expanded nodes**.

We want to find a cheap **path** from Lugoj to Bucharest. Initially, E is **empty**, and F **contains** only Lugoj. We **terminate** if E **contains** Bucharest.

Expansion of a **node** n in F moves it from F to E and adds to F every **neighbor** of n that is not already in E or F . We obtain $g(n)$ by **minimizing** $g(e) + c(e, n)$ over **expanded nodes** e .

As a **heuristic** $h(n)$, we use the **straight-line distance** from n to Bucharest as given by the table in the **lectures**.

Explain how the following **algorithms** choose which **node** to **expand** next:

1. **greedy search** with **heuristic** h

Solution: The **search expands** the **node** n that **minimizes** the **function** $h(n)$.

2. **A* search** with **path cost** g and **heuristic** h

Solution: The **search expands** the **node** n that **minimizes** the **function** $g(n) + h(n)$

3. Explain what h^* is here and why h is **admissible**.

Solution: $h^*(n)$ is the **cost** of the **shortest path** from n to Bucharest (which we do not **know** unless we **expand all nodes**). Because $c(m, n)$ is always bigger than the **straight-line distance**, every **path** is longer than the **straight-line distance** between its end points. Thus $h(n) \leq h^*(n)$.

4. For each `search`, give the `order` in which `nodes` are `expanded`.
(You only have to give the `nodes` to get the full score. But to get partial credit in case you're wrong, you may want to include for each step all `nodes` in the `fringe` and their `cost`.)
-

Solution: Lugoj (244), Mehadia (241), Drobeta (242), Craiova (160), Pitesti (100), Bucharest (0)

astarSearch-search: Lugoj (0+244), Mehadia (70+241), Drobeta ((70+75)+242), Craiova (70+75+120)+160), Timisoara (111+329), Pitesti ((70+75+120+138)+100), Bucharest ((70+75+120+138+101)+0)

Problem 3.4 (Heuristics)

Consider `heuristic search` with `heuristic h`.

1. Briefly `explain` what is the same and what is different between `A*` search and `greedy search` regarding the decision which `node` to `expand` next.
-

Solution: Both choose the `node` that `minimizes` a certain `function`. As that `function`, `A*` uses the `sum of path cost` and `heuristic` whereas `greedy` only uses the `heuristic`.

2. Is the `constant function` $h(n) = 0$ an `admissible heuristic` for `A*` search?
-

Solution: Yes. (But it's a useless one.)

Problem 3.5 (Tree Search in ProLog)

Implement the following tree search algorithms in Prolog:

1. `BFS`
2. `DFS`
3. `Iterative Deepening` (with `step size 1`)

Remarks:

- In the `lectures`, we talked about `expanding nodes`. That is relevant in many `AI applications` where the `tree` is not built yet (and maybe even too big to hold in `memory`), such as `game trees` in `move-based games` or `decision trees` of `agents interacting` with an `environment`. In those cases, when visiting a `node`, we have to `expand` it, i.e., `compute` what its `children` are.

In this problem, we work with smaller **trees** where the **search algorithm** receives the fully **expanded tree** as **input**. The **algorithm** must still visit every **node** and perform some operation on it — the **search algorithm** determines in which **order** the **nodes** are visited.

In our case, the operation will be to *write out the label* of the **node**.

- In the **lectures**, we worked with **goal nodes**, where the **search** stops when a **goal node** is found. Here we do something simpler: we *visit all the nodes and operate on each one* without using a **goal state**. (Having a **goal state** is then just the special case where the operation is to **test** the **node** and possibly stop.)

Concretely, your submission **must** be a single **Prolog file** that extends the following **implementation**:

```
% tree(V,TS) represents a tree.
% V must be a string - the label/value/data V of the root node
% TS must be a list of trees - the children/subtrees of the root node
% In particular, a leaf node is a tree with the empty list of children
istree(tree(V,TS)) :- string(V), istreelist(TS).

% istreelist(TS) holds if TS is a list of trees.
% Note that the children are a list not a set, i.e., they are ordered.
istreelist([]).
istreelist([T|TS]) :- istree(T), istreelist(TS).

% The following predicates define search algorithms that take a tree T
% and visit every node each time writing out the line D:L where
% * D is the depth (starting at 0 for the root)
% * L is the label

% dfs(T) visits every node in depth-first order
dfs(T) :- ????
% bfs(T) visits every node in breadth-first order
bfs(T) :- ????
% itd(T):- visits every node in iterative deepening order
itd(T) :- ????
```

Here “must” means you can define any **number** of additional **predicates**. But the **predicates** specified above must exist and must have that **arity** and must work correctly on any **input** **T** that satisfies **istree(T)**. “Working correctly” means the **predicates** must write out exactly what is specified, e.g.,

```
0:A
1:B
```

for the **depth-first search** of the **tree** **tree("A", [tree("B", [])])**.

Solution:

```
% initialize with depth 0
dfs(T) :- dfsD(T,0).
```

```

% write out depth and value V of the current node,
% then search all children with depth D+1
dfsD(tree(V,TS), D) :- write(D), write(":"), writeln(V),
    Di is D+1, dfsAll(TS,Di).

% calls dfsD on all trees in a list
dfsAll([],_).
dfsAll([T|TS],D) :- dfsD(T,D), dfsAll(TS,D).

% initialize with the fringe containing T at depth 0
 bfs(T) :- bfsFringe([(0,T)]).

% empty fringe - done
bfsFringe([]).

% take the first pair (D,T) in the fringe, write out D and the
% value V of T, append children TS of T paired with depth D+1
% to the *end* of F, and recurse
bfsFringe([(D,tree(V,TS))|F]) :- write(D), write(":"), writeln(V),
    Di is D+1, pair(Di,TS, DTS), append(F, DTS, F2), bfsFringe(F2).

% pair(D,L,DL) takes value D and list L and pairs every element
% in L with D, returning DL
pair(_,[],[]).
pair(D,[H|T],[ (D,H)|DT]) :- pair(D,T,DT).

% initialize with cutoff 0
itd(T) :- itdUntilDone(T,0),!.

% calls dfsUpTo with cutoff C and initial depth
itdUntilDone(T,C) :- dfsUpTo(T,0,C,Done), increaseCutOffIfNotDone(T,C,Done).
% depending on the value of Done, terminate or increase the cutoff.
increaseCutOffIfNotDone(_,_,Done) :- Done=1.
increaseCutOffIfNotDone(T,C,Done) :- Done=0, Ci is C+1, itdUntilDone(T,Ci).

% dfsUpTo(T,D,U,Done) is like dfs(T,D) except that
% * we stop at cutoff depth U
% * we return Done (0 or 1) if there were no more nodes to explore

% cutoff depth reach, more nodes left
dfsUpTo(_, D, U, Done) :- D > U, Done is 0.
% write data, recurse into all children with depth D+1
dfsUpTo(tree(V,TS), D, U, Done) :- write(D), write(":"), writeln(V),
    Di is D+1, dfsUpToAll(TS,Di,U, Done).

% dfsUpToAll(TS,D,U,Done) calls dfsUpTo(T,D,U,_) on all elements of TS;
% it returns 1 if all children did
dfsUpToAll([],_,_,Done) :- Done is 1.
dfsUpToAll([T|TS],D,U,Done) :- dfsUpTo(T,D,U,DoneT),
    dfsUpToAll(TS,D,U,DoneTS), Done is DoneT*DoneTS.

```
